

# Geospatial Fundamentals in R with sf, Part 2

*Patty Frontiera and Drew Hart, UC Berkeley D-Lab*

*August 2019*

## Part II Prep

1. Open the repo at <https://github.com/dlab-berkeley/Geospatial-Fundamentals-in-R-with-sf>
  - Download and unzip the zip file
  - Take note of where the folder is located
2. Start RStudio and open a **new script**, or `./docs/02-spatial_analysis.Rmd`
3. Set your working directory to the folder you unzipped
4. Install the required libraries in RStudio, ONLY IF YOU DO NOT HAVE THEM ALREADY!

```
our_packages<- c("ggplot2", "dplyr", "sf", "units", "tmap")
for (i in our_packages) {
  if ( i %in% rownames(installed.packages()) == FALSE) {
    install.packages(i)
  }
}
```

5. Open the slides, `./docs/02-spatial-analysis.html`, in your browser (or click the “Part 2 Slides” link the repo).

## Part II Overview

Recap Part I

Tour of Spatial Analysis

## Part I Recap

In Part I, we:

- Loaded geospatial data from CSV files
- Mapped data with `ggplot`
- Promoted data frames to `sf` objects with `sf::st_as_sf`
- Loaded geodata from shapefiles with `sf::st_read`
- Explored CRSs with `sf::st_crs`
- Transformed CRSs with `sf::st_transform`
- Mapped data with `tmap`

## R Spatial Libraries

Let's load the libraries we will use

```
library(sf)      # spatial objects and methods
library(tmap)    # mapping spatial objects
```

## Set your working directory

Use `setwd` to set your working directory to the location of the tutorial files.

For example:

```
setwd("~/Documents/Dlab/workshops/2018/rgeo/r-geospatial-workshop/r-geospatial-workshop")
```

## Reload Part I data

You may want to reload the data that we had in our workspace at the end of Part I.

We've provided a little script for doing that, which you can run using the following line of code:

```
source('./docs/reload_part_01_data.R')
```

## Spatial Analysis

### The Spatial Analysis Workflow

1. Mapping / plotting to see location and distribution
2. Asking questions of, or querying, your data
3. Cleaning & reshaping the data
4. Applying analysis methods
5. Mapping analysis results
6. Repeat as needed

### Transform data to common CRS

In order to perform spatial analysis we need to first convert all data objects to a common CRS.

Which type? Projected or Geographic CRS?

### Geographic vs. Projected CRS

If my goal is to create maps, I may convert all data to a geographic CRS.

- Why? Which one?

If my goal is to do spatial analysis, I will convert to a projected CRS.

- Why? Which one?

### Common CRS EPSG Codes

#### Geographic CRSs

- 4326 Geographic, WGS84 (default for lon/lat)
- 4269 Geographic, NAD83 (USA Fed agencies like Census)

#### Projected CRSs

- 5070 USA Contiguous Albers Equal Area Conic
- 3310 CA Albers Equal Area
- 26910 UTM Zone 10, NAD83 (Northern Cal)
- 3857 Web Mercator (web maps)

## Transform all layers to UTM 10N, NAD83

Use `st_transform` to transform `SFhomes15_sf` and `bart` to UTM 10N, NAD83

- `SFhighways` and `SFboundary` already have this CRS

Recall, this transformation is called **projecting** or **reprojecting**

The EPSG code is **26910**, units are meters.

## Transform all layers to UTM 10, NAD83

First, transform `SFhomes15_sf`

(Remember, this is also called *reprojecting*.)

Note the two methods for doing same thing:

```
#highways are already in 26910!
```

```
st_crs(SFhighways)
```

```
## Coordinate Reference System:
```

```
##   EPSG: 26910
```

```
##   proj4string: "+proj=utm +zone=10 +datum=NAD83 +units=m +no_defs"
```

```
#so we can use them as the target CRS
```

```
SFhomes15_utm <- st_transform(SFhomes15_sf, st_crs(SFhighways))
```

```
#OR we could just use the EPSG code directly
```

```
#SFhomes15_utm <- st_transform(SFhomes15_sf, 26910)
```

## Transform the boundary?

```
# Check the CRS
```

```
st_crs(SFboundary) == st_crs(SFhomes15_utm)
```

```
## [1] FALSE
```

```
# Transform
```

```
SFboundary_utm <- st_transform(SFboundary, st_crs(SFhomes15_utm))
```

```
# Check again
```

```
st_crs(SFboundary_utm) == st_crs(SFhomes15_utm)
```

```
## [1] TRUE
```

## BART data - Challenge

Transform the `bart_sf` object to UTM 10N.

Name the new object `bart_utm`

## Challenge: Solution

```
# Transform Bart to UTM
```

```
bart_utm <- st_transform(bart_sf, st_crs(SFhomes15_utm))
```

## Check

Do the CRSs all match?

```
st_crs(bart_utm)$epsg
```

```
## [1] 26910
```

```
st_crs(SFboundary_utm)$epsg
```

```
## [1] 26910
```

```
st_crs(SFhighways)$epsg
```

```
## [1] 26910
```

```
st_crs(SFhomes15_utm)$epsg
```

```
## [1] 26910
```

## Map all layers

Visual check

```
plot(SFboundary_utm)
```

```
lines(SFhighways, col='purple', lwd=4)
```

```
## Error in data.matrix(x): (list) object cannot be coerced to type 'double'
```

```
points(SFhomes15_utm)
```

```
## Warning in data.matrix(x): NAs introduced by coercion
```

```
## Warning in data.matrix(x): NAs introduced by coercion
```




```
## Warning in data.matrix(x): NAs introduced by coercion
```

```
## Warning in data.matrix(x): NAs introduced by coercion
```

```
## Error in data.matrix(x): (list) object cannot be coerced to type 'double'
```

```
plot(bart_utm, col="red", pch=15, add=T)
```

```
## Warning in plot.sf(bart_utm, col = "red", pch = 15, add = T): ignoring all  
## but the first attribute
```



## Map all layers

What happened?


Two things:

1. Remember, by default, `sf`'s `plot` method will plot a grid of maps, one for each variable in the `data.frame`!
2. We can't just plot `sf` objects directly with calls to R's `lines` and `points` functions.

## Map all layers

However, we can get what we want easily, with the help of the `st_geometry` function:

```
plot(st_geometry(SFboundary_utm))
plot(st_geometry(SFhighways), col='purple', lwd=4, add = T)
plot(st_geometry(SFhomes15_utm), add = T, pch = 19, cex = 0.5)
plot(st_geometry(bart_utm), col="skyblue", pch=19, cex = 1, add=T)
```



### Challenge (Optional / time permitting)


Create the same plot, as closely as possible, using `tmap`.

### Challenge: Solution

```
challenge_map = tm_shape(SFboundary) +  
  tm_polygons() +  
tm_shape(SFhighways) +  
  tm_lines(col = 'purple', lwd = 4) +  
tm_shape(SFhomes15_sf) +  
  tm_dots(col = 'black', size = 0.5) +  
tm_shape(bart_utm) +  
  tm_dots(col = 'skyblue', size = 1)
```

### Challenge: Solution

```
challenge_map
```



## Spatial Queries

### Spatial Queries

There are two key types of spatial queries

- **spatial measurement** queries,
  - e.g. area, length, distance
- **spatial relationship** queries,
  - e.g. what locations in A are also in B.

These types are often combined, e.g.

- What is the area of region A that is within region B?

## Spatial Measurement Queries

### Computing Area

What is the area of San Francisco?

What data would we use to answer that question?

### Area of San Francisco

- Use `sf::st_area` to compute the area of `sf` objects with polygons
- Check results against Wikipedia for SF

```
sf_area = st_area(SFboundary_utm)
sf_area
```

```
## 119949901 [m^2]
```

## Area of San Francisco

How did it manage to give us the units?

That comes from the `units` package, which `sf` imports and uses!

```
class(sf_area)
```

```
## [1] "units"
```

```
typeof(sf_area)
```

```
## [1] "double"
```

## Area in sq km

Compare to the Wikipedia page's area for SF

```
sf_area / (1000 * 1000) # Convert to square KM
```

```
## 119.9499 [m^2]
```

## Area in sq km

That number is right, but now we've got an annoying little problem: Our value in square kilometers is labeled as square meters!

The `units` package, an `sf` dependency, provides a better way.

```
library(units)
```

```
## udunits system database from /usr/share/udunits
```

```
set_units(sf_area, km^2)
```

```
## 119.9499 [km^2]
```

## Area in sq km

The function `valid_udunits` will give us a table of the valid units we could convert to:

(Note that the 'ud' comes from the `udunits` package, a dependency of the `units` package.)

```
head(valid_udunits(), 2)
```

```
## udunits system database read from /usr/share/udunits
```

```
## # A tibble: 2 x 11
```

```
##   symbol symbol_aliases name_singular name_singular_a~ name_plural
```

```
##   <chr>   <chr>           <chr>           <chr>           <chr>
```

```
## 1 m      ""             meter          metre          ""
```

```
## 2 kg     ""             kilogram       ""             ""
```

```
## # ... with 6 more variables: name_plural_aliases <chr>, def <chr>,
```

```
## #   definition <chr>, comment <chr>, dimensionless <lgl>, source_xml <chr>
```



## Area of San Francisco

What if we gave `st_area` the SF boundary in an unprojected CRS?

```
st_area(SFboundary)
```

## Area of San Francisco

```
st_area(SFboundary)
```

```
## 120038745 [m^2]
```

`st_area` still gives us the measurement in a reasonable unit (rather than squared decimal degrees).

(However, this isn't a reason not to choose a reasonable, projected CRS for our data! Still best practice.

(Also notice the slight difference in our answers. This is not an equal-area projection!)

```
st_area(SFboundary_utm)
```

```
## 119949901 [m^2]
```

## Length of highways

Use the function `st_length` to compute length of linear geometries.

```
st_length(SFhighways)
```

```
## Units: [m]
```

```
## [1] 106.81285 106.81176 111.51019 111.50913 111.59378 111.59378
## [7] 112.02125 112.01864 110.76642 111.31424 113.21844 113.63364
## [13] 109.70803 109.79960 106.00554 106.00187 106.28392 106.29112
## [19] 105.21542 105.22731 106.39351 106.38778 109.62289 109.62900
## [25] 113.93788 113.93659 54.70773 55.80114 54.71353 55.79882
## [31] 56.86037 55.79639 56.84993 55.80369 56.68567 54.90766
## [37] 56.63758 54.94442 55.49930 56.47812 110.61722 53.84482
## [43] 54.79581 53.63327 54.00083 53.39926 54.32512 53.38106
## [49] 54.17142 54.12614 54.66743 53.40842 53.96184 109.13590
## [55] 54.61882 54.37285 54.84968 54.77273 108.97194 54.33636
## [61] 57.63089 110.55822 208.94243 209.05216 56.59410 53.67826
## [67] 110.31555 108.50992 108.50737 140.98684 54.22959 53.02420
## [73] 34.22720 71.54877 69.91093 69.90407 153.78462 214.54450
## [79] 97.90823 98.13696 241.35294 240.57488 200.78614 93.35164
## [85] 95.85929 83.63511 122.56650 79.05691 91.61146 81.18008
## [91] 97.64588 84.83871 84.53116 138.24165 136.15610 272.46325
## [97] 72.06197 77.57336 78.17869 43.96122 88.45841 79.16623
## [103] 30.10884 37.52394 76.91117 81.12105 53.90307 76.37346
## [109] 34.27302 80.96283 81.27529 51.27864 180.65158 80.83910
## [115] 77.97029 73.19650 219.33740 132.40054 82.67087 181.75832
## [121] 84.60513 95.72339 623.40101 624.13256 1250.25180 764.47517
## [127] 766.94192 160.89800 204.46082 428.26660 113.70954 212.18632
## [133] 103.59201 36.06458 117.74497 21.14108 132.30204 137.46452
## [139] 819.56269 816.21275 211.16053 209.84847 210.83762 210.30138
## [145] 210.89008 210.78982 217.65637 217.59643 216.63497 216.65846
## [151] 211.83011 211.42684 212.15515 209.11608 47.56151 53.09755
## [157] 153.68153 153.34954 153.01433 152.32205 150.89362 150.33399
## [163] 150.61101 150.63365 151.20176 151.20023 150.60030 150.57768
## [169] 150.84118 150.81207 150.91343 150.91135 150.84600 150.89499
```

```
## [175] 150.91640 150.97727 150.88674 150.88337 149.72000 149.73431
## [181] 376.11733 98.16757 154.78278 43.28478 338.17768 96.13831
## [187] 173.47910 195.88397 65.32746 167.10764 312.72440 70.90073
## [193] 71.23036 614.68089 618.12994 369.24483 80.26625 325.36115
## [199] 110.06959 74.52087 79.25213 518.91854 149.83566 231.09698
## [205] 202.09511 237.65820 230.24111 230.32594 385.18077 80.66169
## [211] 156.86509 156.81320 243.26701 243.50972 317.30484 317.20170
## [217] 210.81728 211.14779 210.19899 210.06277 210.59325 209.85169
## [223] 209.37551 209.47373 209.18052 209.08728 209.69002 209.74257
## [229] 209.45700 209.61743 209.42360 209.60218 209.60701 209.56462
## [235] 210.87848 210.80542 211.21313 212.05797 210.47259 210.41286
## [241] 210.54065 210.93940 209.85168 209.11913 212.40005 211.85840
```

## Length of highways

Oh! We got the length of every segment, in meters.

How do we get the total length of highways, in km?

## Challenge

Calculate the total length of SF highways in our dataset, in km.

## Challenge: solution

```
tot_length = set_units(sum(st_length(SFhighways)), km)
tot_length
```

```
## 39.83624 [km]
```

## Perimeter

We can also calculate the perimeter of polygons, should we need it (though this is implemented in the `lwgeom` package, an `sf` dependency, rather than in `sf` itself.)

```
perim = lwgeom::st_perimeter(tracts)
head(perim, 10)
```

```
## Units: [m]
## [1] 3095.519 2771.519 4053.398 2960.833 5846.816 6509.614 2498.040
## [8] 3030.683 3744.900 4150.271
```

## Distance

The `st_distance` will return the min distance between two geometries.

Compute the distance in kilometers between Embarcadero & Powell St Bart stations

(NOTE: You can always spot-check on Google Maps.)

```
emb_pow_dist = st_distance(bart_utm[bart_utm$STATION == 'EMBARCADERO',],
                           bart_utm[bart_utm$STATION == 'POWELL STREET',])
emb_pow_dist = set_units(emb_pow_dist, km)
emb_pow_dist
```

```
## Units: [km]
##          [,1]
## [1,] 1.334997
```

## Distance

Take note of the print-out. What's up with the [1,] and [,1] around the value?

`st_distance` is going to calculate a matrix of pairwise distances, by default! (We just happened to subset our `sf` object to two new objects, each with a single feature, i.e. row.)

Read the docs:

```
?st_distance
```

## Challenge

That means we can easily calculate the distance between all SF properties and Embarcadero station. So go ahead and do that!

### Challenge: solution

```
dist2emb <- st_distance(bart_utm[bart_utm$STATION == 'EMBARCADERO',],
                        SFhomes15_utm)
dist2emb <- set_units(dist2emb, km)

# check output
length(dist2emb)

## [1] 835
nrow(SFhomes15_utm)

## [1] 835
head(dist2emb, 10)

## Units: [km]
## [1]  9.3560500  3.7753035  2.3817736 11.1305421  4.4098830  1.6601826
## [7]  6.9005955  4.2245048 11.1512413  0.6239989
```

### Challenge: solution

Different syntax, equivalent result:

You could just nest your calls, if you'd like.

```
dist2emb <- set_units(st_distance(bart_utm[bart_utm$STATION == 'EMBARCADERO',],
SFhomes15_utm), km)

# check output
head(dist2emb, 10)

## Units: [km]
## [1]  9.3560500  3.7753035  2.3817736 11.1305421  4.4098830  1.6601826
## [7]  6.9005955  4.2245048 11.1512413  0.6239989
```

### Challenge: solution

Different syntax, equivalent result:

You could also use the 'tidy' syntax, if you're into that!

```
dist2emb <- st_distance(bart_utm[bart_utm$STATION == 'EMBARCADERO',],
                        SFhomes15_utm) %>% set_units(km)
# check output
head(dist2emb, 10)

## Units: [km]
## [1]  9.3560500  3.7753035  2.3817736 11.1305421  4.4098830  1.6601826
## [7]  6.9005955  4.2245048 11.1512413  0.6239989
```

## Spatial Relationship Queries

### Spatial Relationship queries

**Spatial relationship queries** compare the geometries of two spatial objects in the same coordinate space (CRS).

Some example relationships:

### Spatial Relationship queries

There are many, often similar, functions to perform spatial relationship queries (can be confusing!).

These operations may return logical values, lists, matrices, dataframes, geometries or spatial objects

- you need to check what type of object is returned
- you need to check what values are returned to make sure they make sense

### BART stations in SF?

This is a very common type of spatial query called a **point-in-polygon** query.

We can use the `st_within` function to answer this.

We'll start with the simplest question: **Are there BART stations in SF?**

We already know the answer, but let's see how it's done.

### Are there any BART stations in SF?

What does it return by default?

```
bart_stations_in_sf <-st_within(bart_utm, SFboundary_utm)

head(bart_stations_in_sf)

## [[1]]
## integer(0)
##
## [[2]]
## integer(0)
##
## [[3]]
## integer(0)
##
## [[4]]
## integer(0)
##
```

```
## [[5]]
## integer(0)
##
## [[6]]
## integer(0)
```

## BART stations in SF?

The docs for the function (`?st_within`) explain that it returns a sparse-matrix object by default. This is more efficient, but more complicated to work with. For our purposes, let's disable this behavior:

```
bart_stations_in_sf <-st_within(bart_utm, SFboundary_utm, sparse=F)

head(bart_stations_in_sf)
```

```
##      [,1]
## [1,] FALSE
## [2,] FALSE
## [3,] FALSE
## [4,] FALSE
## [5,] FALSE
## [6,] FALSE
```

## BART stations in SF?

That's a bit more obvious! Looks like we got a logical value for each BART station.

Let's check the object's size:

```
dim(bart_stations_in_sf)
```

```
## [1] 44  1
```

```
dim(bart_utm)
```

```
## [1] 44  5
```

## BART stations in SF?

So, to answer the simple question, we just need to know if there's at least one TRUE in that list.

```
T %in% bart_stations_in_sf
```

```
## [1] TRUE
```

## Which Bart stations are in SF?

What about this question?

We can use the same output, but now leverage its station-by-station structure!

## Challenge

Return the names of the BART stations that are within SF.

## Challenge: solution

```
bart_utm[bart_stations_in_sf, ]$STATION
```

```
## [1] "EMBARCADERO"          "MONTGOMERY STREET"
## [3] "POWELL STREET"        "CIVIC CENTER/ UN PLAZA"
## [5] "16TH STREET & MISSION" "24TH STREET & MISSION"
## [7] "GLEN PARK"            "BALBOA PARK"
```

## Which Bart stations are in SF?

And of course, there are multiple ways to do a thing!

We could also use the `st_intersection` function to get similar results.

```
sfbart_utm = st_intersection(bart_utm, SFboundary_utm)
```

```
## Warning: attribute variables are assumed to be spatially constant
## throughout all geometries
```

```
sfbart_utm
```

```
## Simple feature collection with 8 features and 7 fields
## geometry type: POINT
## dimension: XY
## bbox: xmin: 548723.7 ymin: 4175071 xmax: 553175.5 ymax: 4183045
## epsg (SRID): 26910
## proj4string: +proj=utm +zone=10 +datum=NAD83 +units=m +no_defs
##          STATION OPERATOR DIST CO          NAME Pop2010 Land_sqmi
## 30          EMBARCADERO      BART    4 SF San Francisco 805235    46.87
## 31      MONTGOMERY STREET      BART    4 SF San Francisco 805235    46.87
## 32          POWELL STREET      BART    4 SF San Francisco 805235    46.87
## 33 CIVIC CENTER/ UN PLAZA      BART    4 SF San Francisco 805235    46.87
## 34  16TH STREET & MISSION      BART    4 SF San Francisco 805235    46.87
## 35  24TH STREET & MISSION      BART    4 SF San Francisco 805235    46.87
## 36              GLEN PARK      BART    4 SF San Francisco 805235    46.87
## 37          BALBOA PARK      BART    4 SF San Francisco 805235    46.87
##          geometry
## 30 POINT (553175.5 4183045)
## 31 POINT (552693.7 4182561)
## 32 POINT (552231.4 4182101)
## 33 POINT (551587.2 4181453)
## 34 POINT (551132.8 4179866)
## 35 POINT (551242.5 4178546)
## 36 POINT (549876.6 4176357)
## 37 POINT (548723.7 4175071)
```

## Map the SF BART stations

```
tmap_mode("view")

tm_shape(SFboundary_utm) +
  tm_polygons(col="beige", border.col="black") +
tm_shape(sfbart_utm) +
  tm_dots(col="red")
```

## Map the SF BART stations

### Reset tmap to plot mode

```
tmap_mode("plot")
```

```
## tmap mode set to plotting
```

### st\_within vs st\_intersects vs st\_intersection

#### Devil in the details...

`st_within` returns TRUE/FALSE, testing if one geometry is *completely* within another.

`st_intersects` returns TRUE/FALSE, testing if two geometries have any points in common.

`st_intersection` returns the geometry that intersects.

### st\_within, st\_intersects, st\_intersection, and friends

- These were just a couple examples of common geometric queries used in spatial analysis.
- These, and other similar operations are neatly summarized on this great `sf` cheatsheet (also available in the `./docs` subdirectory of our workshop repo):

## SF Census Tracts

Let's consider the `SFhomes15_utm` data along with the *SF census tract* data that we saw on day 1.

However, we are going to work with another version of the tract data, one that includes the population for each tract.

## Challenge

Read in the SF Census Tracts with pop data and call it `sftracts`

- The filename is `sftracts_wpop.shp`.
- The file is located in `./data`.

Then, create a population choropleth map.

## Challenge: solution

```
#read in tracts
```

```
sftracts <- st_read("./data", "sftracts_wpop")
```

```
## Reading layer `sftracts_wpop' from data source `/home/drew/Desktop/stuff/berk/dlab/Geospatial-Fundam
```

```
## Simple feature collection with 195 features and 10 fields
```

```
## geometry type:  MULTIPOLYGON
```

```
## dimension:      XY
```

```
## bbox:           xmin: -122.5145 ymin: 37.70813 xmax: -122.328 ymax: 37.86334
```


```
## epsg (SRID):    4269
```

```
## proj4string:    +proj=longlat +datum=NAD83 +no_defs
```

## Challenge: solution

```
#plot
```

```
plot(sftracts['pop14'])
```



## Composite operations

### Composite operations

The remaining material will work through some common spatial analysis tasks.

Each workflow will feature some combination spatial measurement operations, spatial relationship operations, and other non-spatial operations.

## Joins and Aggregation

### Spatial join

A spatial join associates rows of data in one object with rows in another object based on the spatial relationship between the two objects.

A spatial join is based on the comparison of two sets of geometries in the same coordinate space.

- This is also called a **spatial overlay**.

### Spatial join

We could use any of a family of spatial relationships that all return matrices of logical values.

`sf` refers to these as ‘geometric binary predicates’, and collects all their documentation into one document, which we’ve already seen:

```
?st_within
```

## In what census tract is each property located?

We need to **spatially join** the `sftracts` and `SFhomes15_utm` to answer this.



What spatial object are we joining data from? to?

## Spatial join

We have points, which are pretty much certain to be either inside or outside polygons. So we'll use `st_within` again as our spatial relationship.

We want to associate with each home the name of the census tract within which it falls.

So here goes...

*In what census tract is each SF property located?*

```
homes_with_tracts <- st_within(SFhomes15_utm, sftracts)
```

## Did it work?

If not, why not?

## CRSs must be the same

The `st_within` function, like almost all spatial analysis functions, requires that both data sets be spatial objects (they are) with the same coordinate reference system (CRS). Let's investigate

```
# What is the CRS of the property data?
st_crs(SFhomes15_utm)

# What is the CRS of the census tracts?
st_crs(sftracts)
```

## Transform the CRS

```
#transform to UTM
sftracts_utm = st_transform(sftracts, st_crs(SFhomes15_utm))

# make sure the CRSs are the same
st_crs(sftracts_utm) == st_crs(SFhomes15_utm)
```

```
## [1] TRUE
```

Now let's try that overlay operation again

## Try 2

*In what tract is each SF property is located?*

```
homes_with_tracts <- st_within(SFhomes15_utm, sftracts_utm)
```

## Review the `st_within` output

What is our output? Does it answer our question?

What type of data object did the over function return?

```
homes_with_tracts <- st_within(SFhomes15_utm, sftracts_utm)

class(homes_with_tracts)
length(homes_with_tracts)
```

```
nrow(sftracts_utm)
nrow(SFhomes15_utm)
```

## Review the `st_within` output

What do we have here?

```
homes_with_tracts <- st_within(SFhomes15_utm, sftracts_utm)
class(homes_with_tracts)
```

```
## [1] "sgbp"
```

```
length(homes_with_tracts)
```

```
## [1] 835
```

```
nrow(sftracts_utm)
```

```
## [1] 195
```

```
nrow(SFhomes15_utm)
```

```
## [1] 835
```

## Review the `st_within` output

What the heck is an object of the class `sgbp`?

**Read the docs!**

(It's basically just a special sparse-matrix structure designed to hold the results returned from these binary-predicate functions.)

```
?sgbp
```

## Review the `st_within` output

What data does the output object *store*?

```
head(homes_with_tracts)
```

```
## [[1]]
## [1] 92
##
## [[2]]
## [1] 38
##
## [[3]]
## [1] 193
##
## [[4]]
## [1] 34
##
## [[5]]
## [1] 153
##
## [[6]]
## [1] 64
```

## Review the `st_within` output

We have a `list`, where each item's *index* is a `SFhomes15_utm` property's index, and each *value* is the index of the `sftracts_utm` census tract within which it is found.

We're halfway there!

## Spatial join

We can now finish the operation by:

1. using that `st_within` output object to subset the `sftracts_utm` `data.frame`;
2. grabbing the desired columns from that subsetted `data.frame` and adding them to our `SFhomes15_utm` `data.frame`.

In our case, the desired column will just be the `GEOID` column (a standardized ID that we can then use to link up to non-spatial census data).

## Add the `GEOID` column

*CAUTION: this only works because the data are in the right order!*

```
SFhomes15_utm$home_geoid <- sftracts_utm[unlist(homes_with_tracts),]$GEOID
```

## Check the result

```
head(SFhomes15_utm, 2)
```

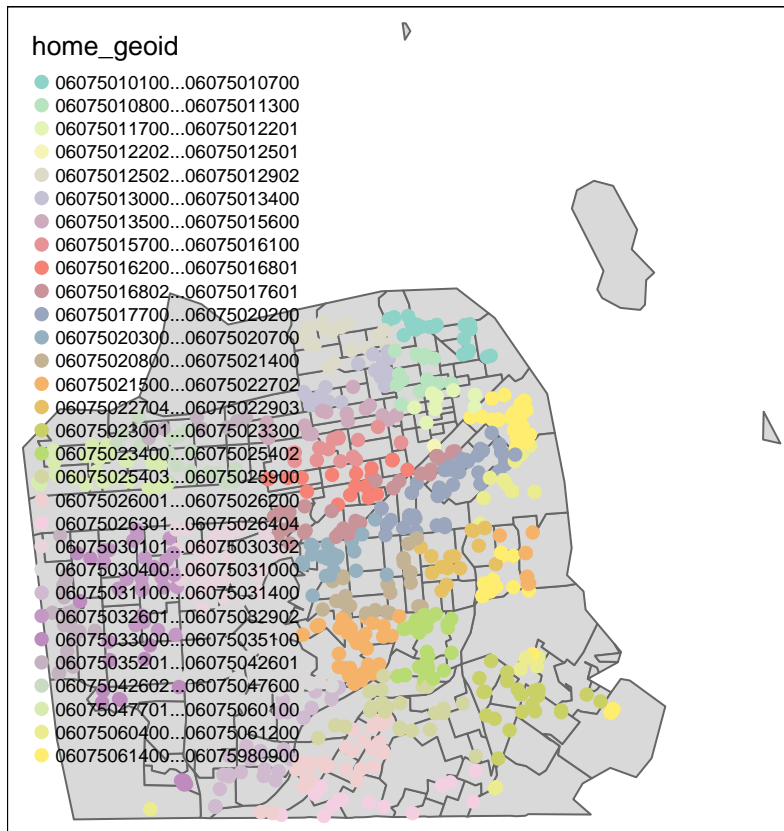
```
## Simple feature collection with 2 features and 18 fields
## geometry type:  POINT
## dimension:      XY
## bbox:           xmin: 546340.7 ymin: 4176656 xmax: 553157.3 ymax: 4179269
## epsg (SRID):    26910
## proj4string:     +proj=utm +zone=10 +datum=NAD83 +units=m +no_defs
##   FiscalYear SalesDate Address YearBuilt
## 24      2015 2015-08-21 0000 2760 19TH      AV0015      1979
## 35      2015 2015-08-13 0000 0560AMISSOURI ST0000      2003
##   NumBedrooms NumBathrooms NumRooms NumStories NumUnits AreaSquareFeet
## 24           2           2         5           0         1           1595
## 35           2           2         5           1         1           1191
##   ImprovementValue LandValue Neighborhood
## 24          432500    432500 West of Twin Peaks
## 35          701280    701280 Potrero Hill
##                                     Location SupeDistrict totvalue SalesYear
## 24 (37.7360097396496, -122.474067310226)          7    865000      2015
## 35 (37.759197817252, -122.396516184449)         10   1402560      2015
##                                     geometry home_geoid
## 24 POINT (546340.7 4176656) 06075030800
## 35 POINT (553157.3 4179269) 06075061400
```

## Check the result

```
join_map = tm_shape(sftracts_utm) +
  tm_polygons() +
  tm_shape(SFhomes15_utm) +
  tm_dots(col = 'home_geoid', size = 0.25)
```

## Check the result

*#Note that tmap bins our tracts because we have so many*  
join\_map



## WOW

Data linkage via space!

The `st_within` operation gave us the census tract data info for each point in `SFhomes15_utm`

We added the GEOID for each point to the `SFhomes15_utm` sf object.

We can now join `SFhomes15_utm` points by GEOID to any census variable, eg median household income, and then do an analysis of the relationship between, for example, property value and that variable.

**How would we do that?**

## Attribute Joins

### Attribute Joins

Attribute joins merge data in two tables based on matching data values contained in a column in each table.

For example we could join a table of student grades with a table of student names and addresses if both tables contain a column with student id.

## Read in the census data

Let's read in a CSV file of median household income for SF tracts.

The `sf_med_hh_income2015.csv` file only has two columns: `GEOID` and `medhhinc`.

Because `GEOIDs` can have leading zeros, we set the `colClasses` to make sure they are not stripped.

```
med_hh_inc <- read.csv("data/sf_med_hh_income2015.csv", stringsAsFactors = F,  
                      colClasses = c("character", "numeric"))
```

```
head(med_hh_inc)
```

```
##           GEOID medhhinc  
## 1 06075980401         0  
## 2 06075990100         0  
## 3 06075012502       11925  
## 4 06075012301       13909  
## 5 06075061100       16545  
## 6 06075980501       16638
```

## Joining a regular data.frame to an sf data.frame

We can use `merge` to join the `med_hh_inc` DF to the `SFhomes15_utm` sf object.

We should make sure that they share a column of common values - `GEOID` / `home_geoid`

## Joining a regular data.frame to an sf data.frame

Join two data objects based on common values in a column.

Use `merge` to join two data.frames.

(Notice, again, that our `sf data.frame` will conveniently behave just like regular old `data.frame` in this way.)

```
#make sure we're using `base` `merge` (because multiple other packages  
#that you might have read in also have a `merge` function)  
SFhomes15_utm <- base::merge(SFhomes15_utm,  
                             med_hh_inc, by.x="home_geoid", by.y="GEOID")
```

## Take a look at output

```
head(SFhomes15_utm, 2) # Look for the col medhhinc
```

```
## Simple feature collection with 2 features and 19 fields  
## geometry type:  POINT  
## dimension:      XY  
## bbox:           xmin: 551575.9 ymin: 4184223 xmax: 551672.2 ymax: 4184228  
## epsg (SRID):    26910  
## proj4string:     +proj=utm +zone=10 +datum=NAD83 +units=m +no_defs  
##   home_geoid FiscalYear SalesDate Address  
## 1 06075010100      2015 2015-06-04 0000 0650 CHESTNUT      ST0204  
## 2 06075010100      2015 2015-01-08 0592 0588 CHESTNUT      ST0000  
##   YearBuilt NumBedrooms NumBathrooms NumRooms NumStories NumUnits  
## 1      1995           2           2         0           0         1  
## 2      1907           0           0        17           3         1  
##   AreaSquareFeet ImprovementValue LandValue Neighborhood
```

```
## 1      1103      571078      571078 North Beach
## 2      3264      654836      1527951 North Beach
##                               Location SupeDistrict totvalue SalesYear
## 1 (37.8039342366397, -122.41411670973)      3  1142156      2015
## 2 (37.8039733892367, -122.413021836652)      3  2182787      2015
## medhhinc      geometry
## 1      61442 POINT (551575.9 4184223)
## 2      61442 POINT (551672.2 4184228)
```

## Check the merge results

```
tmap_mode("view")
tm_shape(SFhomes15_utm) + tm_dots(col="medhhinc")
```

## The Census Tract Perspective

We now know the census tract for each property.

Now let's think about this question from the tract perspective.

Let's ask the question

- What is the average property value per tract?

## Non-Spatial Aggregation

Since we joined GEOID to each property we can use the non-spatial `aggregate` function to compute the mean of `totvalues` for each GEOID.

But we'll use `sf`'s spatial implementation of `aggregate`.

We'll start by...

Reading the docs!

```
?sf::aggregate.sf
```

### `sf::aggregate.sf`

We see that we can provide arguments:

- **x**: `sf` object to be aggregated
- **by**: can be another `sf` object whose geometries will generate the groupings
- **FUN**: function to be used to summarize the grouped values

## What is the mean home value in each census tract?

```
tracts_with_mean_val <- aggregate(x = SFhomes15_utm["totvalue"],
                                by = sftracts_utm,
                                FUN = mean)
```

Wow, so simple. What does that give us?

## Examine output of `sf::aggregate.sf`

```
class(tracts_with_mean_val)
```

```
## [1] "sf"          "data.frame"
head(tracts_with_mean_val, 2)

## Simple feature collection with 2 features and 1 field
## geometry type:  MULTIPOLYGON
## dimension:      XY
## bbox:          xmin: 551221.8 ymin: 4182036 xmax: 552338.1 ymax: 4184030
## epsg (SRID):   26910
## proj4string:    +proj=utm +zone=10 +datum=NAD83 +units=m +no_defs
##   totvalue      geometry
## 1      NA MULTIPOLYGON (((551683.5 41...
## 2  482000 MULTIPOLYGON (((551221.8 41...

nrow(tracts_with_mean_val) == nrow(sftracts_utm)

## [1] TRUE
```

## sf::aggregate.sf output

`sf::aggregate.sf` returned a new `sf data.frame`.

The new `data.frame` has the same geometry as `sftracts_utm`

But it only contains one column, with the mean `totvalue` for each tract.

To make these data more useful, let's add this value to `sftracts_utm`!

**Note:** This only works because there are the same number of elements in both `data.frames` and they are in the same order!

```
sftracts_utm$mean_totvalue <- tracts_with_mean_val$totvalue

head(sftracts_utm, 2) # check it

## Simple feature collection with 2 features and 11 fields
## geometry type:  MULTIPOLYGON
## dimension:      XY
## bbox:          xmin: 551221.8 ymin: 4182036 xmax: 552338.1 ymax: 4184030
## epsg (SRID):   26910
## proj4string:    +proj=utm +zone=10 +datum=NAD83 +units=m +no_defs
##   STATEFP COUNTYFP TRACTCE      AFFGEOID      GEOID   NAME LSAD
## 1      06      075  010700 14000000US06075010700 06075010700   107   CT
## 2      06      075  012201 14000000US06075012201 06075012201 122.01  CT
##   ALAND AWATER pop14      geometry mean_totvalue
## 1 183170      0  5311 MULTIPOLYGON (((551683.5 41...      NA
## 2  92048      0  4576 MULTIPOLYGON (((551221.8 41...  482000
```

## Map it

Map the results to make sure they seem reasonable.

(NOTE: This is called a ‘choropleth’ map.)

```
choropleth =
tm_shape(sftracts_utm) +
  tm_polygons(col="mean_totvalue", border.col=NA)
```

## Map it

```
choropleth
```

## Why no values for some tracts?

```
choropleth + tm_shape(SFhomes15_utm) + tm_dots(size = 0.01)
```

## Proximity Analysis

Many methods of spatial analysis use distance to select features. For example...

*What properties are within walking distance of BART?*

In order to select properties with 1KM of BART, we can:

1. create a 1km-radius buffer polygon around each BART point
2. do a point-in-polygon operation to either count the number of properties within the buffer or compute mean values.

## Create the buffers

For this, we'll use—surprise, surprise—`st_buffer`.

But first, we'll...

Read the docs!

```
?st_buffer
```

## Create the buffers

It takes as input:

- **x**: an `sf*` object or objects to be buffered;
- **dist**: a buffer distance.

## Create the buffers

Let's assume 1km is our 'standard walking distance'.

```
#remember: our units are meters!  
bart_1km_buffer <- st_buffer(sfbart_utm, dist=1000)
```

## Map the buffers

```
tm_shape(bart_1km_buffer) + tm_polygons(col="red") +  
tm_shape(sfbart_utm) + tm_dots()
```

## What properties are within 1km of a bart station?

What operation can we use here?

Once again, we can use `st_intersects` or `st_intersection`



## What properties are within 1km of a bart station?

```
SFhomes_near_bart <-st_intersection(SFhomes15_utm, bart_1km_buffer)
```

```
## Warning: attribute variables are assumed to be spatially constant
## throughout all geometries
```

```
# Take a look
```

```
head(SFhomes_near_bart)
```

```
## Simple feature collection with 6 features and 26 fields
```

```
## geometry type: POINT
```

```
## dimension: XY
```

```
## bbox: xmin: 552402.6 ymin: 4182488 xmax: 552895.5 ymax: 4183658
```

```
## epsg (SRID): 26910
```

```
## proj4string: +proj=utm +zone=10 +datum=NAD83 +units=m +no_defs
```

```
##   home_geoid FiscalYear SalesDate Address
## 19 06075010500      2015 2015-08-14 0000 0016 FRONT ST0000
## 21 06075010500      2015 2015-09-08 0000 0733 FRONT ST0707
## 24 06075010600      2015 2015-12-16 0000 0455 VALLEJO ST0311
## 42 06075011700      2015 2015-04-17 0000 0690 MARKET ST1905
## 43 06075011700      2015 2015-06-05 0000 0333 BUSH ST3804
## 44 06075011700      2015 2015-03-12 0000 0333 BUSH ST3904
```

```
##   YearBuilt NumBedrooms NumBathrooms NumRooms NumStories NumUnits
## 19      1986           3           3         7           2         1
## 21      2007           2           2         5           1         1
## 24      1973           0           0         3           0         1
## 42      2007           1           1         4           0         1
## 43      1987           2           2         5           0         1
## 44      1987           2           2         5           2         1
```

```
##   AreaSquareFeet ImprovementValue LandValue
## 19          2251          1475000  1475000
## 21          1378          1200000  1200000
## 24           825           437500   437500
## 42           952           477167   715751
## 43          1510           812200   812200
## 44          1510           761361   761361
```

```
##   Neighborhood Location
## 19 Financial District/South Beach (37.7982533719796, -122.399172133445)
## 21 Financial District/South Beach (37.7980267794367, -122.400091849915)
## 24 North Beach (37.7987906647799, -122.404766465894)
## 42 Financial District/South Beach (37.7882423441494, -122.403200587421)
## 43 Financial District/South Beach (37.7905798690843, -122.403108701388)
## 44 Financial District/South Beach (37.7905798690843, -122.403108701388)
```

```
##   SupeDistrict totvalue SalesYear medhhinc STATION OPERATOR DIST CO
## 19           3  2950000      2015   105000 EMBARCADERO BART 4 SF
## 21           3  2400000      2015   105000 EMBARCADERO BART 4 SF
## 24           3   875000      2015   34808 EMBARCADERO BART 4 SF
## 42           3  1192918      2015   34914 EMBARCADERO BART 4 SF
## 43           3  1624400      2015   34914 EMBARCADERO BART 4 SF
## 44           3  1522722      2015   34914 EMBARCADERO BART 4 SF
```

```
##   NAME Pop2010 Land_sqmi geometry
## 19 San Francisco 805235 46.87 POINT (552895.5 4183601)
## 21 San Francisco 805235 46.87 POINT (552814.7 4183576)
## 24 San Francisco 805235 46.87 POINT (552402.6 4183658)
```

```
## 42 San Francisco 805235 46.87 POINT (552547.9 4182488)
## 43 San Francisco 805235 46.87 POINT (552554.4 4182748)
## 44 San Francisco 805235 46.87 POINT (552554.4 4182748)
```

## Plot it

```
tmap_mode('view')

## tmap mode set to interactive viewing

tm_shape(bart_1km_buffer) + tm_borders(col="red") +
tm_shape(sfbart_utm) + tm_dots() +
tm_shape(SFhomes_near_bart) +
tm_dots(col = 'green', size = 0.03)
```

## Any Questions?

### Summary

That was a whirlwind tour of just some of the methods of spatial analysis.

There was of course a lot we didn't and can't cover.

### Selected References & Tutorials

Here's that great `sf` cheatsheet (also available in the `./docs` subdirectory of this repo).

Introductory tutorials

- Spatial Data in R tutorial
- NEON Spatial Data tutorials
- GIS in R

### Selected References & Tutorials

Emphasis on geodata visualization

- Tmap in a Nutshell
- Intro to visualizing Spatial Data in R
- RStudio Leaflet in R tutorial
- Blog on mapping census data in R

### Selected references & tutorials

Deep dive Tutorials that include spatial analysis

- Geocomputation in R
- Intro to GIS and Spatial Analysis (see appendices)
- An Introduction to Spatial Data Analysis and Visualisation in R

CRAN Spatial Packages

- CRAN Task View: Analysis of Spatial Data